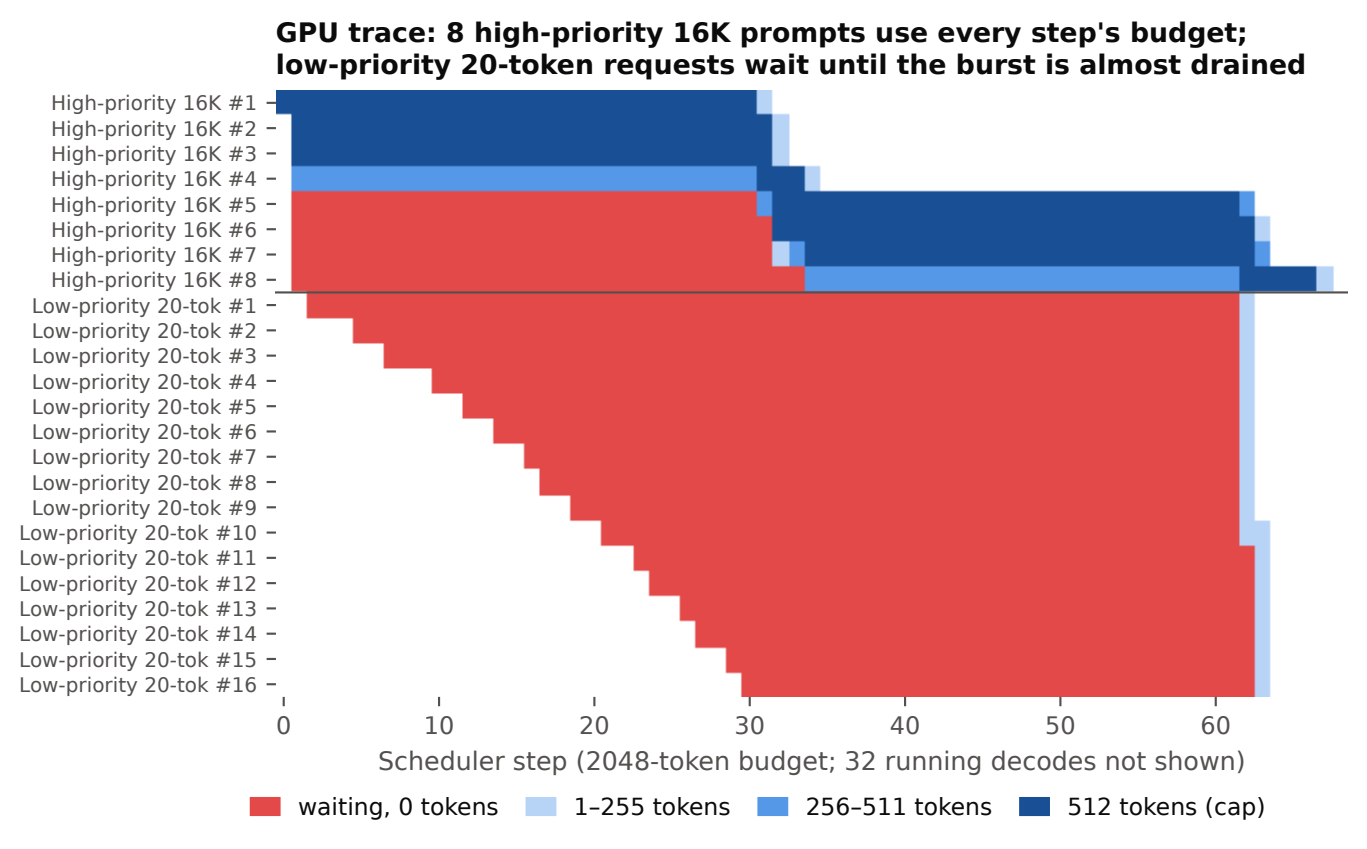
GPU trace: 8 high-priority 16K prompts use every step's budget; low-priority 20-token requests wait until the burst is almost drained
High-priority 16K #1
High-priority 16K #2
High-priority 16K #3
High-priority 16K #4
High-priority 16K #5
High-priority 16K #6
High-priority 16K #7
High-priority 16K #8
Low-priority 20-tok #1
Low-priority 20-tok #2
Low-priority 20-tok #3
Low-priority 20-tok #4
Low-priority 20-tok #5
Low-priority 20-tok #6
Low-priority 20-tok #7
Low-priority 20-tok #8
Low-priority 20-tok #9
Low-priority 20-tok #10
Low-priority 20-tok #11
Low-priority 20-tok #12
Low-priority 20-tok #13
Low-priority 20-tok #14
Low-priority 20-tok #15
Low-priority 20-tok #16
0
10
20
30
40
50
60
Scheduler step (2048-token budget; 32 running decodes not shown)
waiting, 0 tokens
1–255 tokens
256–511 tokens
512 tokens (cap)