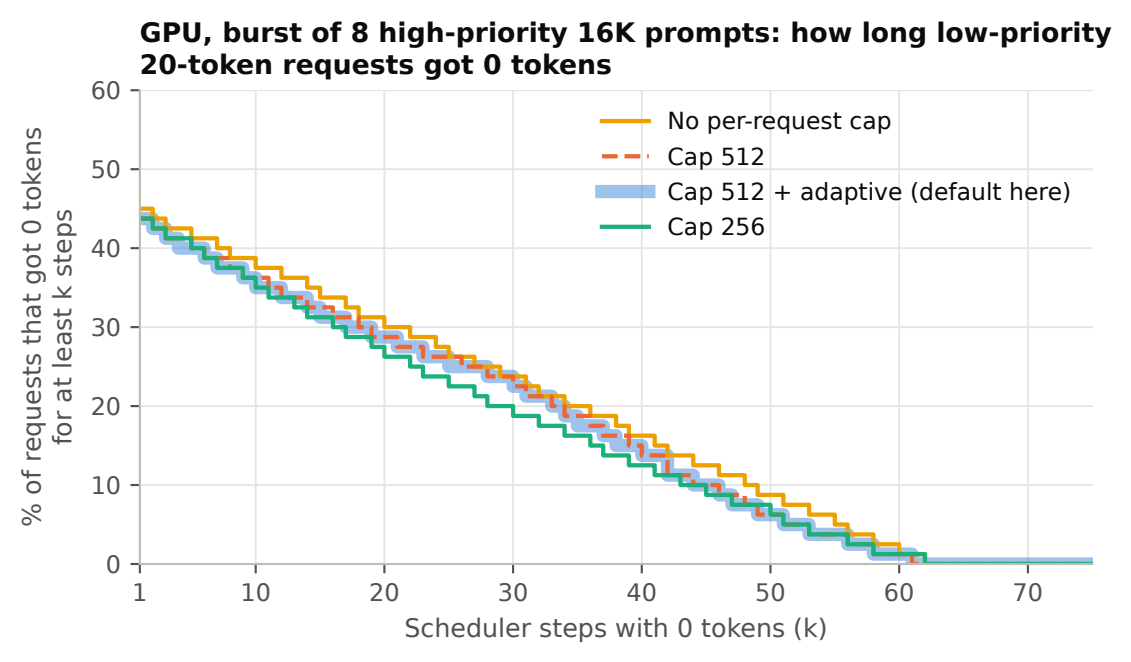
GPU, burst of 8 high-priority 16K prompts: how long low-priority 20-token requests got 0 tokens
No per-request cap
Cap 512
Cap 512 + adaptive (default here)
Cap 256
% of requests that got 0 tokens for at least k steps
Scheduler steps with 0 tokens (k)
0
10
20
30
40
50
60
1
10
20
30
40
50
60
70