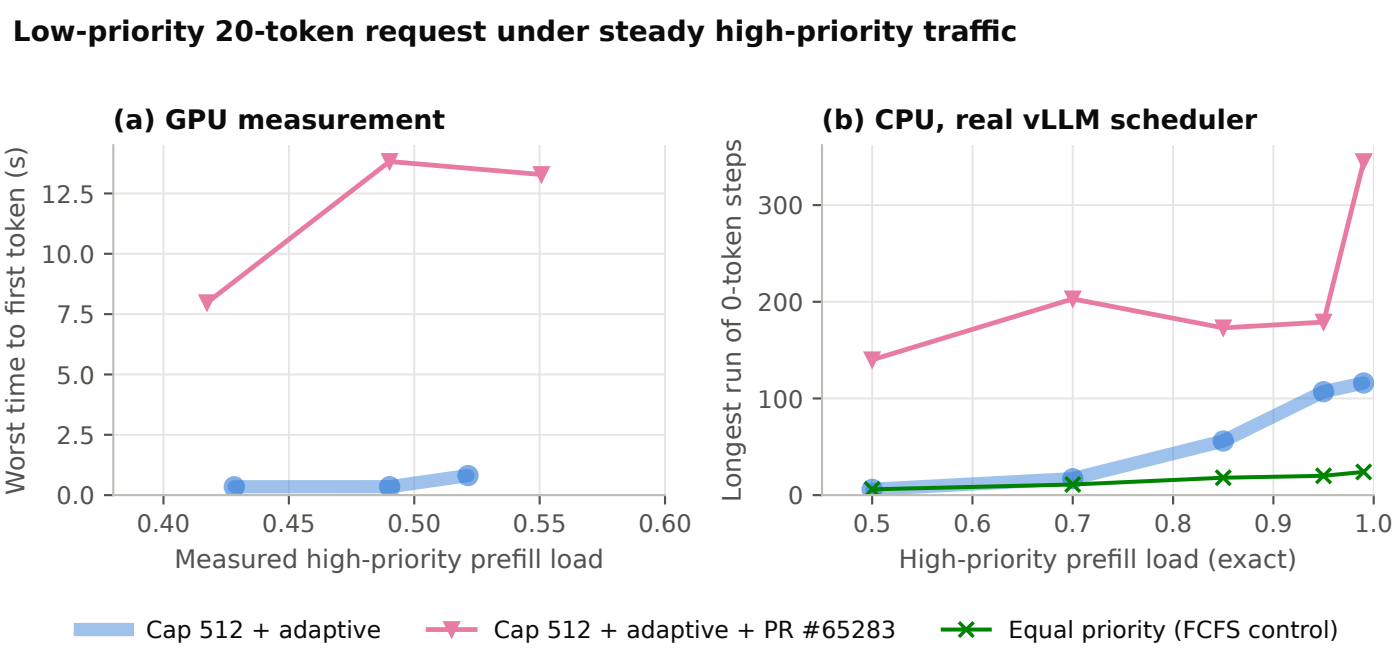
Low-priority 20-token request under steady high-priority traffic
(a) GPU measurement
Worst time to first token (s)
0.0
2.5
5.0
7.5
10.0
12.5
0.40
0.45
0.50
0.55
0.60
Measured high-priority prefill load
(b) CPU, real vLLM scheduler
Longest run of 0-token steps
0
100
200
300
0.5
0.6
0.7
0.8
0.9
1.0
High-priority prefill load (exact)
Cap 512 + adaptive
Cap 512 + adaptive + PR #65283
Equal priority (FCFS control)